

50.021 – AI

Alex

Week 06: Markov Decision Processes II

[The following notes are compiled from various sources such as textbooks, lecture materials, Web resources and are shared for academic purposes only, intended for use by students registered for a specific course. In the interest of brevity, every source is not cited. The compiler of these notes gratefully acknowledges all such sources.]

Key takeaways:

Be able to explain the main ideas behind:

- explain the difference between equations for optimal policy (uses $\max_a T(a)$) and for a given policy (uses $\sum_a T(a)\pi(a|s)$)
- explain the difference between fixpoint computations for optimal policy (uses $\max_a T(a)$) and for a given policy (uses $\sum_a T(a)\pi(a|s)$)
- Q-function $Q^\pi(s, a)$
- Q-iteration to get the Q-function $Q^{\pi^*}(s, a)$ for the optimal policy π^*

1 Recap

MDP recap

an MDP is a 4-tuple: S, A, P, R

- have world state space S , action space A
- $P(s'|a, s)$ – transition probability when taking in state s action a to arrive in state s'
- $R(s, a, s')$ reward function (can be a probability e.g. $R(s, a, s') = N(\hat{R}(s, a, s'), \sigma^2)$, where $\hat{R}(s, a, s')$ is a function)

policy

Policy $\pi(a|s)$ – what action to choose given that one is in state s . Can be a probability over all allowed actions.

Policy is just a rule (deterministic as $a = \pi(s)$ or stochastic with $\sum_a \pi(a|s) = 1$) what action a to choose in state s .

What is the performance measure of an agent? **Expected (future) reward**

for unbounded time horizons and in general with a discount factor $\gamma \in (0, 1]$

$$\bar{r}(s) = E_{(\dots, a_t \sim \pi(a|s_t=s_t), \dots)} \left[\sum_{t=0}^{\infty} \gamma^t r_t | s_0 = s \right]$$

the expectation runs over drawing an action for every time step – future reward is a function of the policy π !

the Value function

The value $V^\pi(s)$ of a policy π is the gamma-discounted expected future reward when starting in state s and continuing according to the policy:

$$V^\pi(s) = E_{(a_0 \sim \pi(a|s_0=s), a_1 \sim \pi(a|s_1=s_1), \dots, a_t \sim \pi(a|s_t=s_t), \dots)} \left[\sum_{t=0}^{T_e} \gamma^t r_t | s_0 = s \right]$$

A policy π^* is optimal if

$$\forall s \quad V^{\pi^*}(s) = \sup_{\pi} V^\pi(s)$$

the goal in MDPs

goal: find a policy π^* which maximizes V^π , the expected reward, averaged over all uncertainties in state transitions, actions and rewards.

The Bellman equation

If rewards and policy are deterministic, that is when the policy π returns one state $a = \pi(s)$ as a function of the current state, then

$$V^\pi(s) = \sum_{s'} P(s'|s, a = \pi(s)) R(s, a = \pi(s), s') + \gamma \sum_{s'} P(s'|s, a = \pi(s)) V^\pi(s')$$

In the general case, replacing any terms depending on an action $T(a = \pi(s)) \rightsquigarrow \sum_a \pi(a|s) T(a)$ by the expectation under probability of an action (, we have:

$$V^\pi(s) = \sum_{s'} \sum_a \pi(a|s) P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} \sum_a \pi(a|s) P(s'|s, a) V^\pi(s')$$

This applies for $T(a) = R(s, a, s')$ and for $T(a) = P(s'|a, s)$

1.1 The Bellman optimality criterion

The Bellman optimality criterion

In an MDP the optimal Value and the optimal policy satisfy the following equations:

$$V^{\pi^*}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s')$$

$$\pi^*(s) = \operatorname{argmax}_{s'} \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s')$$

1.2 Estimating $V^{\pi^*}(s)$ and $\pi^*(s)$

Ok, we know what condition they should satisfy, ...

By fixpoint iteration! Start with $V_0(s) = 0$. Iterate: Compute $V_{k+1}(s)$ from $V_k(s)$

$$V^{\pi^*}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s')$$

$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V_k(s')$$

until an iteration k such that: $\max_s |V_{k+1}(s) - V_k(s)| < \delta$

2 Comparison to avoid confusion: want Value of the optimal policy π^* or of a given policy π ?

If you want to estimate the value V^{π^*} under **the optimal policy** π^* , then you use

$$V^{\pi^*}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s')$$

$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V_k(s')$$

If you want the Value V^π under a fixed policy π , then the following Bellman equations hold

$$V^\pi(s) = \sum_{s'} P(s'|s, a = \pi(s)) R(s, a = \pi(s), s') + \gamma \sum_{s'} P(s'|s, a = \pi(s)) V^\pi(s')$$

for a deterministic policy $\pi(s)$. Or for a stochastic policy:

$$V^\pi(s) = \sum_{s'} \sum_a P(s'|s, a) R(s, a, s') \pi(a|s) + \gamma \sum_{s'} \sum_a P(s'|s, a = \pi(s)) \pi(a|s) V^\pi(s')$$

As a consequence of that: If you want the Value under a fixed policy π , then this can also be computed by a fixpoint problem but not using \max_a :

$$V^\pi(s) = \sum_{s'} P(s'|s, a = \pi(s)) R(s, a = \pi(s), s') + \gamma \sum_{s'} P(s'|s, a = \pi(s)) V^\pi(s')$$

$$V_{k+1}(s) = \sum_{s'} P(s'|s, a = \pi(s)) R(s, a = \pi(s), s') + \gamma \sum_{s'} P(s'|s, a = \pi(s)) V_k(s')$$

for a stochastic policy:

$$V^\pi(s) = \sum_{s'} \sum_a P(s'|s, a) R(s, a, s') \pi(a|s) + \gamma \sum_{s'} \sum_a P(s'|s, a = \pi(s)) \pi(a|s) V^\pi(s')$$

$$V_{k+1}(s) = \sum_{s'} \sum_a P(s'|s, a) R(s, a, s') \pi(a|s) + \gamma \sum_{s'} \sum_a P(s'|s, a = \pi(s)) \pi(a|s) V_k(s')$$

This converges analogously, as the proof for the Value function under the optimal policy.

Comparison

Value $V^{\pi^*}(s)$ of optimal policy π^*

$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V_k(s')$$

Value $V^\pi(s)$ of any given policy π

$$\text{determ: } V_{k+1}(s) = \sum_{s'} P(s'|s, a = \pi(s)) R(s, a = \pi(s), s') + \gamma \sum_{s'} P(s'|s, a = \pi(s)) V_k(s')$$

$$\text{stoch: } V_{k+1}(s) = \sum_{s'} \sum_a P(s'|s, a) R(s, a, s') \pi(a|s) + \gamma \sum_{s'} \sum_a P(s'|s, a = \pi(s)) \pi(a|s) V_k(s')$$

The difference: for the optimal policy you use $\max_a T(a)$, for a given policy use $T(a = \pi(s))$ or $\sum_a T(a) \pi(a|s)$

3 Q-function

idea of Q-function

Idea: $Q^\pi(s, a)$ is the expected future reward when using action a in state s (not following any policy!), and after that continuing with policy π . This means as a definition:

$$Q^\pi(s, a) = \underbrace{\sum_{s'} P(s'|s, a) R(s, a, s')}_{\text{Reward for doing } a \text{ in state } s} + \underbrace{\gamma \sum_{s'} P(s'|s, a) V^\pi(s')}_{\text{future } E[\cdot] \text{ reward under } \pi \text{ after doing } a \text{ in state } s}$$

Since Value V^π means the expected reward when choosing the action according to the current policy π , the following must hold:

plugging in the policy in Q must result in V :

Q^π -function and Value-function V^π :

$$V^\pi(s) = Q^\pi(s, a = \pi(s))$$

$$V^\pi(s) = \sum_a \pi(a|s) Q^\pi(s, a)$$

This can be plugged in into above definition to yield a definition of Q^π as a fixpoint same as we had for V^π

The Bellman equations for the Q -function:

deterministic policy:

$$Q^\pi(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^\pi(s')$$

$$= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) Q^\pi(s', a' = \pi(s'))$$

stochastic policy:

$$Q^\pi(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^\pi(s')$$

$$= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \sum_{a'} \pi(a'|s') Q^\pi(s', a')$$

Optimal policy means that we choose the optimal action...

Q^{π^*} -function and Value-function V^{π^*} for the optimal policy π^* :

By Bellman-optimality criterion (and the definition of Q^{π^*}) we have **only for the optimal policy π^*** :

$$V^{\pi^*}(s) = \max_a Q^{\pi^*}(s, a)$$

This one can plug in:

The Bellman optimality for the Q -function and the optimal policy π^* :

$$\begin{aligned} Q^{\pi^*}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a') \\ V^{\pi^*}(s) &= \max_a Q^{\pi^*}(s, a) \end{aligned}$$

The difference is the same as for values: for the optimal policy you use $\max_a T(a)$, for a given policy use $T(a = \pi(s))$ or $\sum_a T(a) \pi(a|s)$

Q^π is often easier to analyze than V^π

Q-iteration

Computes the Q -function under the optimal policy π^* .

$$\begin{aligned} Q_{k+1}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a') \\ V^{\pi^*}(s) &= \max_a Q^{\pi^*}(s, a) \\ \pi^*(s) &= \operatorname{argmax}_a Q^{\pi^*}(s, a) \end{aligned}$$

Compare:

Bellman-optimality criterion for Value V^{π^*} under optimal policy π^*

$$V^{\pi^*}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s')$$

Bellman-optimality criterion for Q -function Q^{π^*} under optimal policy π^*

$$Q^{\pi^*}(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a')$$

and here the iteration algorithms:

Value-iteration

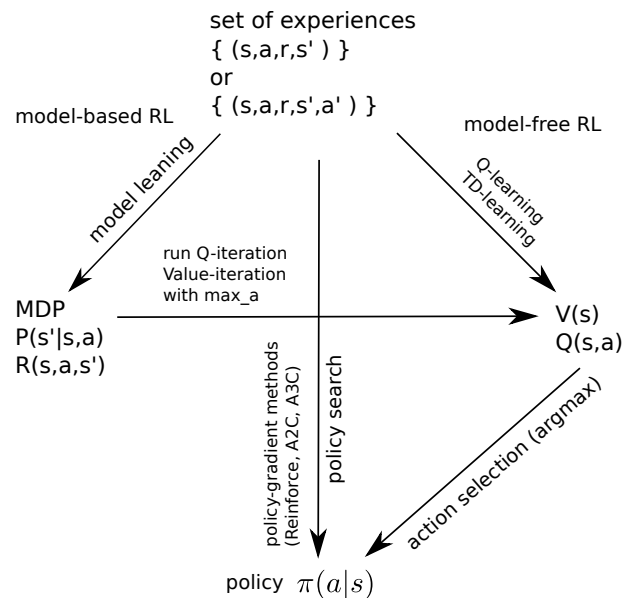
$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V_k(s')$$

Q-iteration

$$\begin{aligned} Q_{k+1}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a') \\ V^{\pi^*}(s) &= \max_a Q^{\pi^*}(s, a) \end{aligned}$$

Not the same.

4 Overview over some types of reinforcement learning



model-free or model-based RL?

Historically two alternative views were developed: cognitive psychologists like Wolfgang Koehler, Edward Tolman (https://en.wikipedia.org/wiki/Wolfgang_K%C3%B6hler https://en.wikipedia.org/wiki/Edward_C._Tolman) on the one side vs the behaviorism of e.g. Edward Thorndike, Clark Hull, perfected by Burrhus Skinner https://en.wikipedia.org/wiki/Edward_Thorndike https://en.wikipedia.org/wiki/Clark_L._Hull https://en.wikipedia.org/wiki/B._F._Skinner

Is learning in animals based on them learning an insight of cause and effect, which can be seen as a model of the process

http://wkprc.eva.mpg.de/english/files/wolfgang_koehler.htm

or did they learn by adapting to stimulus-response scenarios (reinforcement – positive reinforcement by giving positive events, negative reinforcements by removing negative events, e.g. Bart vs Cupcakes) ?

RL carries a lot of ideology from behaviourism.

Draw-backs of a model-free learning approach:

- no generalizable / reusable knowledge, only values for the current goal
- no human-like expectation about future outcomes (what states), only value

- goal changes – learning starts from scratch

Both schools explain some aspects of learning. Touching a hot plate ...

5 the start of Q-learning

Goal: Learn Q-function $Q^{\pi^*}(s, a)$ for optimal policy π^*

Remember:

$$Q^{\pi^*}(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a')$$

This translates to:

Q = expected reward + $\gamma \times Q$ for new states s' averaged with probability to land in s' .

What is when we have observed an experience (s, a, r, s') and want to learn from it?

We can replace expected reward by r . We can replace Q for new states s' averaged ... by Q in our observed state s' , so

$$Q_{\text{current}}(s, a) \approx r + \gamma \max_{a'} Q_{\text{current}}(s', a')$$

in an iterative fashion:

$$s, a \rightsquigarrow r, s' \\ Q_{\text{new}}(s, a) \approx r + \gamma \max_{a'} Q_{\text{old}}(s', a')$$

However we want a slow update, so weight it with α

$$\begin{aligned} Q_{\text{new}}(s, a) &= (1 - \alpha) Q_{\text{old}}(s, a) + \alpha \left(r + \gamma \max_{a'} Q_{\text{old}}(s', a') \right) \\ &= Q_{\text{old}}(s, a) + \alpha \left(r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a) \right) \end{aligned}$$

What is the idea?

$$\begin{aligned} r + \gamma \max_{a'} Q_{\text{old}}(s', a') &> Q_{\text{old}}(s, a) \Rightarrow \text{increase } Q_{\text{old}}(s, a) \\ r + \gamma \max_{a'} Q_{\text{old}}(s', a') &< Q_{\text{old}}(s, a) \Rightarrow \text{decrease } Q_{\text{old}}(s, a) \end{aligned}$$

Compare observed reward r plus $\gamma \times$ your estimated Q from the observed new state s' (from experience (s, a, r, s')) against your estimated $Q(s, a)$ from the observed old state s and update iteratively.

Reinforcement:

- more reward then currently estimated \Rightarrow increase $Q_{\text{old}}(s, a)$
- less reward then currently estimated \Rightarrow decrease $Q_{\text{old}}(s, a)$

Analogously can be executed for the value function.

Q-Learning by TD(0)-Learning for the Q-function

$$\begin{aligned} Q_{\text{new}}(s, a) &= (1 - \alpha)Q_{\text{old}}(s, a) + \alpha \left(r + \gamma \max_{a'} Q_{\text{old}}(s', a') \right) \\ &= Q_{\text{old}}(s, a) + \alpha \left(r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a) \right) \end{aligned}$$

As algorithm?

- init $Q(s, a) = 0$, choose start state s
- run while loop:
 - choose action $a \approx_{\epsilon} \operatorname{argmax}_a Q(s, a)$
 - observe reward r and new state s' to obtain (s, a, r, s')
 - update $Q(s, a) = Q_{\text{old}}(s, a) + \alpha (r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a))$
 - set oldstate to new state $s = s'$

$a \approx_{\epsilon} \operatorname{argmax}_a Q(s, a)$ is what ? – so called ϵ -greedy exploration


$$a = \begin{cases} \text{random} & \text{with prob } \epsilon \\ \operatorname{argmax}_a Q(s, a) & \text{else} \end{cases}$$

Idea: Do not follow strictly your current estimate of best action. Allow a random action with some probability to explore new options.

Limitation: above works for discrete states s . Not if states are continuous.

5.1 Towards Q-learning for continuous states with neural nets and the like

Suppose your states are continuous. Then we can replace a tabular $Q(s, a)$ by a parametrized function $Q_w(s, a)$. Q_w can be e.g. an Multi-layer perceptron with as many outputs as $|A|$, or the like.



How to optimize in that case? We need a loss to compute a gradient.
Can start from:

$$\begin{aligned} s, a &\rightsquigarrow r, s' \\ Q_{\text{new}}(s, a) &\approx r + \gamma \max_{a'} Q_{\text{old}}(s', a') \\ \Rightarrow 0 &\approx r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{new}}(s, a) \end{aligned}$$

This looks like a regression problem, and it can be approached (this is a simplification by)

$$L((s, a, r, s')) = (r + \gamma \max_{a'} Q_w(s', a') - Q_w(s, a))^2$$

as a loss for a single experience (s, a, r, s') . Now can do minibatch-gradient and SGD training or anything similar.