

# 50.021 – AI

Alex

## Week 06: Q-learning I

[The following notes are compiled from various sources such as textbooks, lecture materials, Web resources and are shared for academic purposes only, intended for use by students registered for a specific course. In the interest of brevity, every source is not cited. The compiler of these notes gratefully acknowledges all such sources. ]

### Key takeaways:

Be able to explain the main ideas behind:

- explain basic Q-learning
- explain SARSA
- explain the difference between SARSA and Q-learning
- name three tricks used in the nature paper for atari games
- explain the idea behind n-step bootstrap methods

## 1 Q-function

### idea of Q-function

Idea:  $Q^\pi(s, a)$  is the expected future reward when using action  $a$  in state  $s$  (not following any policy!), and after that continuing with policy  $\pi$ . This means as a definition:

The Bellman equations for the  $Q$ -function:

deterministic policy:

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^\pi(s') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) Q^\pi(s', a' = \pi(s')) \end{aligned}$$

stochastic policy:

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^\pi(s') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \sum_{a'} \pi(a'|s') Q^\pi(s', a') \end{aligned}$$

The Bellman optimality for the  $Q$ -function and the optimal policy  $\pi^*$ :

$$\begin{aligned} Q^{\pi^*}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) V^{\pi^*}(s') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a') \\ V^{\pi^*}(s) &= \max_a Q^{\pi^*}(s, a) \end{aligned}$$

The difference is the same as for values: for the optimal policy you use  $\max_a T(a)$ , for a given policy use  $T(a = \pi(s))$  or  $\sum_a T(a) \pi(a|s)$   
Q-iteration: Computes the  $Q$ -function under the optimal policy  $\pi^*$ .

$$\begin{aligned} Q_{k+1}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q_k(s', a') \\ V^{\pi^*}(s) &= \max_a Q^{\pi^*}(s, a) \\ \pi^*(s) &= \operatorname{argmax}_a Q^{\pi^*}(s, a) \end{aligned}$$

## 2 the start of Q-learning

Goal: Learn  $Q$ -function  $Q^{\pi^*}(s, a)$  for optimal policy  $\pi^*$

Remember:

$$Q^{\pi^*}(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a')$$

This translates to:

$Q$  = expected reward +  $\gamma \times Q$  for new states  $s'$  averaged with probability to land in  $s'$ .

What is when we have observed an experience  $(s, a, r, s')$  and want to learn from it?

Note:  $a \sim \pi(a|s)$ ,  $s' \sim P(s'|a, s)$ ,  $r = R(s, a, s')$

We can approximate expected reward by  $r$ :

$$\sum_{s'} P(s'|s, a) R(s, a, s') \approx r$$

We can replace  $Q$  for new states  $s'$  averaged ... by  $Q$  in our observed state  $s'$ :

$$\begin{aligned} \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a') &\approx \gamma \max_{a'} Q_{\text{current}}(s', a'), \text{ so:} \\ Q^{\pi^*}(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^{\pi^*}(s', a') \\ Q_{\text{current}}(s, a) &\approx r + \gamma \max_{a'} Q_{\text{current}}(s', a') \end{aligned}$$

in an iterative fashion:

$$\begin{aligned} s, a &\rightsquigarrow r, s' \\ Q_{\text{new}}(s, a) &\approx r + \gamma \max_{a'} Q_{\text{old}}(s', a') \end{aligned}$$

However we want a slow update, so weight it with  $\alpha$

$$\begin{aligned} Q_{\text{new}}(s, a) &= (1 - \alpha) Q_{\text{old}}(s, a) + \alpha \left( r + \gamma \max_{a'} Q_{\text{old}}(s', a') \right) \\ &= Q_{\text{old}}(s, a) + \alpha \left( r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a) \right) \end{aligned}$$

What is the idea?

$$\begin{aligned} r + \gamma \max_{a'} Q_{\text{old}}(s', a') &> Q_{\text{old}}(s, a) \Rightarrow \text{increase } Q_{\text{old}}(s, a) \\ r + \gamma \max_{a'} Q_{\text{old}}(s', a') &< Q_{\text{old}}(s, a) \Rightarrow \text{decrease } Q_{\text{old}}(s, a) \end{aligned}$$

Compare observed reward  $r$  plus  $\gamma \times$  your estimated  $Q$  from the observed new state  $s'$  (from experience  $(s, a, r, s')$ ) against your estimated  $Q(s, a)$  from the observed old state  $s$  and update iteratively.

Reinforcement:

- more reward then currently estimated  $\Rightarrow$  increase  $Q_{\text{old}}(s, a)$
- less reward then currently estimated  $\Rightarrow$  decrease  $Q_{\text{old}}(s, a)$

Analogously can be executed for the value function.

#### Q-Learning by TD(0)-Learning for the Q-function

$$\begin{aligned} Q_{\text{new}}(s, a) &= (1 - \alpha)Q_{\text{old}}(s, a) + \alpha \left( r + \gamma \max_{a'} Q_{\text{old}}(s', a') \right) \\ &= Q_{\text{old}}(s, a) + \alpha \left( r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a) \right) \end{aligned}$$

As algorithm?

- init  $Q(s, a) = 0$ , choose start state  $s$
- run while loop:
  - choose action  $a \approx_{\epsilon} \arg\max_a Q(s, a)$
  - observe reward  $r$  and new state  $s'$  to obtain  $(s, a, r, s')$
  - update  $Q(s, a) = Q_{\text{old}}(s, a) + \alpha (r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a))$
  - set oldstate to new state  $s = s'$

$a \approx_{\epsilon} \arg\max_a Q(s, a)$  is what ? – so called  $\epsilon$ -greedy exploration

$$a = \begin{cases} \text{random} & \text{with prob } \epsilon \\ \arg\max_a Q(s, a) & \text{else} \end{cases}$$

Idea: Do not follow strictly your current estimate of best action. Allow a random action with some probability to explore new options.

Limitation: above works for discrete states  $s$ . Not if states are continuous.

Q-learning as above finds the  $Q^{\pi^*}$  for the optimal policy  $\pi^*$

### 3 SARSA: Q-function estimation for a given policy

above:  $Q^{\pi^*}$ . Here goal  $Q^{\pi}$  for a given policy  $\pi$ .

Remember:

$$Q^{\pi}(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) Q^{\pi}(s', a' = \pi(s'))$$

When we have an experience  $s, a, r, s'$ , then  $s' \sim P(s'|s, a)$ . If we assume also that  $a' = \pi(s')$ , then:

$$\begin{aligned} \sum_{s'} P(s'|s, a) R(s, a, s') &\approx r \\ \gamma \sum_{s'} P(s'|s, a) Q^{\pi}(s', a' = \pi(s')) &\approx \gamma Q^{\pi}(s', a') \end{aligned}$$

Therefore we can use  $r + \gamma Q^\pi(s', a')$  to update  $Q^\pi(s, a)$  softly:

$$\begin{aligned} Q_{\text{new}}^\pi(s, a) &= (1 - \alpha)Q_{\text{old}}^\pi(s, a) + \alpha(r + \gamma Q_{\text{old}}^\pi(s', a')) \\ &= Q_{\text{old}}^\pi(s, a) + (r + \gamma Q_{\text{old}}^\pi(s', a') - Q_{\text{old}}^\pi(s, a)) \end{aligned}$$

The difference between Q-learning and SARSA is: SARSA is a so-called on-policy method.

**on-policy is defined as computing quantities for the given policy  $\pi$ .** While off policy denotes the optimization for a different policy (usually some greedy policy or an  $\epsilon$ -greedy policy.) There are two variants: SARSA for policy evaluation and SARSA for policy control (= learning the optimal policy  $\pi^*$ ).

#### SARSA for policy evaluation with a given policy $\pi$

goal  $Q^\pi$  for a given policy  $\pi$ .


Given an experience  $s, a = \pi(s), r, s', a' = \pi(s')$ ,  
initialize  $\forall(s, a) \ Q_{\text{new}}^\pi(s, a) = 0$ , then iterate until convergence:

- update  $Q$

$$Q_{\text{new}}^\pi(s, a) = Q_{\text{old}}^\pi(s, a) + (r + \gamma Q_{\text{old}}^\pi(s', a') - Q_{\text{old}}^\pi(s, a))$$

- update  $s = s', a = a'$

### 3.1 When one can use SARSA to optimize for the optimal policy?



SARSA does not try to learn the optimal policy if you choose  $a \sim \pi(a|s)$  for a fixed  $\pi$ !

How to learn the optimal policy with sarsa? One needs to combine sarsa with iteratively changing the policy towards the greedy policy. This can be done by replacing  $a \sim \pi(a|s)$  by doing in every step:

$$a \approx_{\epsilon} \operatorname{argmax}_a Q_t(s, a) \text{ and} \\ a' \approx_{\epsilon} \operatorname{argmax}_a Q_t(s', a)$$

When you combine SARSA with choosing  $a \approx_{\epsilon} \operatorname{argmax}_a Q_t(s, a)$  and  $a' \approx_{\epsilon} \operatorname{argmax}_a Q_t(s', a)$ , then it will converge to the optimal policy provided that  $\epsilon = \epsilon(t) \xrightarrow{t} 0$  and that all state-action-pairs are visited an infinite number of times.

### SARSA for policy control – obtaining $\pi^*$

goal  $Q^{\pi^*}$  for the optimal policy  $\pi^*$ .

Given an experience  $s, a \approx_{\epsilon} \operatorname{argmax}_a Q_t(s, a), r, s', a' \approx_{\epsilon} \operatorname{argmax}_a Q_t(s', a')$ ,  
 initialize  $\forall(s, a) Q_{\text{new}}(s, a) = 0$ , then iterate until convergence:

- update  $Q$

$$Q_{\text{new}}(s, a) = Q_{\text{old}}(s, a) + (r + \gamma Q_{\text{old}}(s', a') - Q_{\text{old}}(s, a))$$


- update  $s = s', a = a'$

See algorithm on page 106 for the version of sarsa which converges to the optimal policy.

### 3.2 When SARSA obtains a different result from Q-Learning while using $\epsilon$ -greedy policy?

The famous example is the cliffworld on page 108. Q-learning assumes using the optimal policy. Whenever the  $\epsilon$ -greedy policy has significant risks, and thus is a lot different from following the optimal policy, then SARSA will give a different solution.

## 4 Towards Deep Q-learning for continuous states with neural nets and the like



How to optimize in that case? We need a loss to compute a gradient.  
 Can start from:

$$\begin{aligned} s, a &\rightsquigarrow r, s' \\ Q_{\text{new}}(s, a) &\approx r + \gamma \max_{a'} Q_{\text{old}}(s', a') \\ &\Rightarrow 0 \approx r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{new}}(s, a) \end{aligned}$$

This looks like a regression problem, and it can be approached (this is a simplification by)

$$L((s, a, r, s')) = (r + \gamma \max_{a'} Q_w(s', a') - Q_w(s, a))^2$$

as a loss for a single experience  $(s, a, r, s')$ . Now can do minibatch-gradient and SGD training or anything similar.

Mnih et al. Nature paper <https://www.nature.com/articles/nature14236> has many modifications – all for the goal to make learning more stable.

Please take a look at this paper.

## 4.1 Fix Q-Target

Consider:

$$L((s, a, r, s')) = (r + \gamma \max_{a'} Q_w(s', a') - Q_w(s, a))^2$$

Using the same  $Q_w$  for computing the expected reward and for the current  $Q_w(s, a)$  was found to be unstable, a deeper look shows that one should use two  $Q_w$  functions:  $Q_{\text{old}}$  – the target  $Q$ , and  $Q_{\text{new}}$  the  $Q$  for optimization

$$0 \approx (r + \gamma \max_{a'} Q_{\text{old}}(s', a') - Q_{\text{new}}(s, a))^2$$

The first change is to consider

$$w = \operatorname{argmin}_w (r + \gamma \max_{a'} Q_{\text{target}}(s', a') - Q_w(s, a))^2$$

where we keep  $Q_{\text{target}}$  constant for  $K$  episodes. Only every  $K$  episodes we update it by

$$Q_{\text{target}} = Q_w \text{ if } n_{\text{episode}} \% K == 0$$

We fix the target for a while to avoid the situation that rapid changes in  $Q_w$  (due to the SGD Gradient updates) make the learning unstable.

There is an alternative to that which also often works is to run at every iteration for a small  $\beta \approx 0$ :

$$Q_{\text{target}} = (1 - \beta) * Q_{\text{target}} + \beta Q_w$$

which updates the target only a bit every episode.

## 4.2 Robust losses

The Huber loss:

$$l(y, z) = 0.5(y - z)^2 1[|y - z| < 1] + |y - z| 1[|y - z| \geq 1]$$

has smaller gradients for large deviations than the MSE loss.

Also: clip gradients to lie in  $[-1, +1]$



### 4.3 Experience Replay

$Q$ -learning is an off-policy methods: it optimizes for a different policy than the one used to generate samples.

As such it is possible to generate a stack of  $N$  experiences  $(s, a, r, s')$ . For most games, they use  $N = 10^6$ . At training time one does the following:

- play an episode, insert all experiences  $(s, a, r, s')$  into the memory
- randomly draw one  $(s, a, r, s')$  for gradient descent


### 4.4 Image preprocessing

the Atari output has shape  $(3, 210, 160)$  per frame. They do the following:

- covert to gray, reshape to 84 size. Results in  $(1, 84, 84)$

This is not sufficient.

Trick is: collect every fourth frame one frame, and assemble 4 of this into one input of shape  $(4, 84, 84)$  – so in total collect 4 frames out of 16 into one single input. Guess why not using a single frame?



### 4.5 use a CNN for $Q$

Standard CNN architecure to process  $(4, 84, 84)$  inputs.

## 5 n-step bootstrap methods

Idea:

$$Q^\pi(s, a) = \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q^\pi(s', a')$$

is approximated by

$$Q_{\text{new}}(s, a) \approx r + \gamma \max_{a'} Q_{\text{old}}(s', a')$$

In this approximation we take the reward of one step  $r$  and an estimate for the next step. In the beginning of the learning, the estimate  $Q_{\text{old}}(s', a')$  can be very poor, this will inject noise into the update.

Key idea: use more observed rewards before resorting to estimated  $Q_{\text{old}}$ !

Suppose one does in state  $s$  action  $a$ , and then in state  $s'$  action  $a'$ , and then continues according to the best  $Q$  in state  $s''$ ?

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) Q^\pi(s', a') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s''|s', a') R(s', a', s'') + \gamma^2 \sum_{s'} P(s''|s', a') R(s', a', s'') \max_{a''} Q^\pi(s'', a'') \end{aligned}$$

One can do that for three steps:

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s'|s, a) Q^\pi(s', a') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s''|s', a') R(s', a', s'') + \gamma^2 \sum_{s'} P(s''|s', a') \max_{a''} Q^\pi(s'', a'') \\ &= \sum_{s'} P(s'|s, a) R(s, a, s') + \gamma \sum_{s'} P(s''|s', a') R(s', a', s'') + \gamma^2 \sum_{s''} P(s'''|s'', a'') R(s'', a'', s''') + \gamma^3 \sum_{s''} P(s'''|s'', a'') \max_{a'''} Q^\pi(s''', a''') \end{aligned}$$

The structure is

$$Q = \text{reward}(\text{step 1}) + \gamma^1 \text{reward}(\text{step 2}) + \gamma^2 \text{reward}(\text{step 3}) + \gamma^3 \sum_{s''} P(s'''|s'', a'') \max_{a'''} Q^\pi(s''', a''')$$

Suppose one observes a chain:  $s_0, a_0, r_0, s_1, a_1, r_1, s_2, a_2, r_2, s_3$

This could be approximated by

$$\begin{aligned} Q_{\text{new}}(s_0, a_0) &= (1 - \alpha) Q_{\text{old}}(s_0, a_0) + \alpha \left( r_0 + \gamma^1 \max_{a_1} Q_{\text{old}}(s_1, a_1) \right) \\ &= (1 - \alpha) Q_{\text{old}}(s_0, a_0) + \alpha \left( r_0 + \gamma^1 r_1 + \gamma^2 \max_{a_2} Q_{\text{old}}(s_2, a_2) \right) \\ &= (1 - \alpha) Q_{\text{old}}(s_0, a_0) + \alpha \left( r_0 + \gamma^1 r_1 + \gamma^2 r_2 + \gamma^3 \max_{a_3} Q_{\text{old}}(s_3, a_3) \right) \end{aligned}$$

This is the basis for so called n-step algorithms such as n-step SARSA.

Advantage: use more real experiences and less to be learned estimates  $Q$ .

For an algorithm: see algorithm on page 120

## 6 n-step bootstrap methods for off-policy learning

This is off graded knowledge. Useful when using replay memory or  $\epsilon$ -greedy policies for sampling.

see page 121 and algorithm on page 122. This is for the case when the policy used for sampling (e.g. from a  $Q$  from a past iteration, as when using replay memory!) differs from the policy  $\pi$  currently being optimized ... which is usually  $\pi(s) = \operatorname{argmax}_a Q(s, a)$  for the current step.